

Executable exact query plan: actual populations shrink before precision escalates

GPU HBM

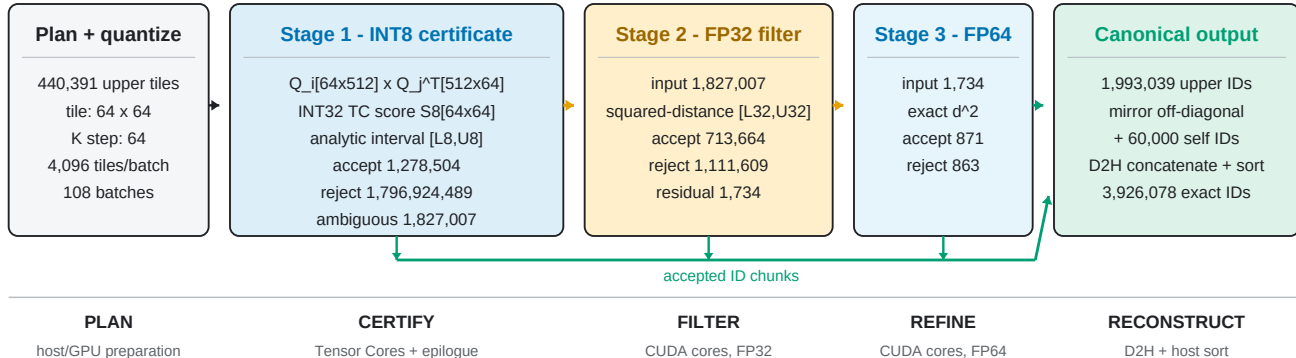
X: 60,000 x 512 FP32

Q: INT8 + per-vector scale/error

fixed-capacity buffers persist across 108 batches

ID buffers + six counters

host count reads set the two dynamic launch extents



All populations are measured in the exact G5 P1 full-scale run; compilation is outside the public timer.

Only unresolved IDs advance; direct accepts bypass later precision stages and enter canonical reconstruction.